

Elena Rostova

Lead Machine Learning Engineer

Email: [elena.rostova@ai-innovations.io](mailto:elena.rostova@ai-innovations.io) | Phone: +44 20 7946 0991 | London, UK

GitHub: <https://github.com/elena-ai>

## SUMMARY

Staff ML Engineer specializing in transformer fine-tuning, embedding retrieval models, and low-latency LLM inference.

## EXPERIENCE

DeepAI Labs | Principal ML Engineer (2020-01 to Present)

- Deployed open-source Gemma-3-4B inference server on serverless GPUs.
- Optimized vector search pipeline with custom HNSW indexing.

## EDUCATION

University of Cambridge | Ph.D. in Computer Science (2015-09 to 2019-06)

## SKILLS

PyTorch, Transformers, SentenceTransformers, vLLM, Python, CUDA, pgvector

## PUBLICATIONS

Optimizing LLM Inference on Serverless GPUs, NeurIPS Workshop 2023